

A controlled evaluation


Wavelet KAN (153K parameters)
vs. four baselines

MIT-BIH, inter-patient
(de Chazal DS1/DS2)

Same protocol and budget
20 paired seeds per model

Validation-based selection
Holm-corrected paired tests

Supraventricular recall is low
for all models (PC-WavKAN 0.20)



Interpretability: not supported

Untrained, prior-centred 0.97

Trained, prior-centred 0.78

Untrained, isotropic 0.67

Trained, isotropic 0.58

Retention of the prior values (PPR)

Training moves randomly initialised
wavelets away from the priors